

Name: Navpreet Kloy  
Wednesday Lab Partner: Robert Armenta

**Model:** Audio-Text-to-Text

**Process:**

1st) Install necessary files:

Python installation check:

```
[MacBook-Air:~ navpreetkloy$ python3 --version  
Python 3.14.2
```

Git verification:

```
[MacBook-Air:~ navpreetkloy$ git --version  
git version 2.50.1 (Apple Git-155)
```

Homebrew installed:

- For easy installation of ffmpeg

```
[MacBook-Air:~ navpreetkloy$ brew --version  
Homebrew 5.1.8
```

FFmpeg installed to process audio files:

```
[MacBook-Air:~ navpreetkloy$ ffmpeg -version  
ffmpeg version 8.1 Copyright (c) 2000-2026 the FFmpeg developers  
built with Apple clang version 21.0.0 (clang-2100.0.123.102)  
configuration: --prefix=/opt/homebrew/Cellar/ffmpeg/8.1_1 --enable-shared --enable-pthreads --enable-version3
```

PyTorch:

```
[MacBook-Air:~ navpreetkloy$ python3 -c "import torch; print(torch.__version__)"  
2.11.0
```

Transformers:

```
[MacBook-Air:~ navpreetkloy$ python3 -c "import transformers; print(transformers.__version__)"  
5.7.0
```

The transcribe.py file was created with code that includes timestamps. These timestamps are in seconds:

```
transcribe.py (3.14.2)
from transformers import pipeline
import sys

if len(sys.argv) < 2:
    print("Usage: python3 transcribe.py <audio_file>")
    sys.exit(1)

audio_file = sys.argv[1]

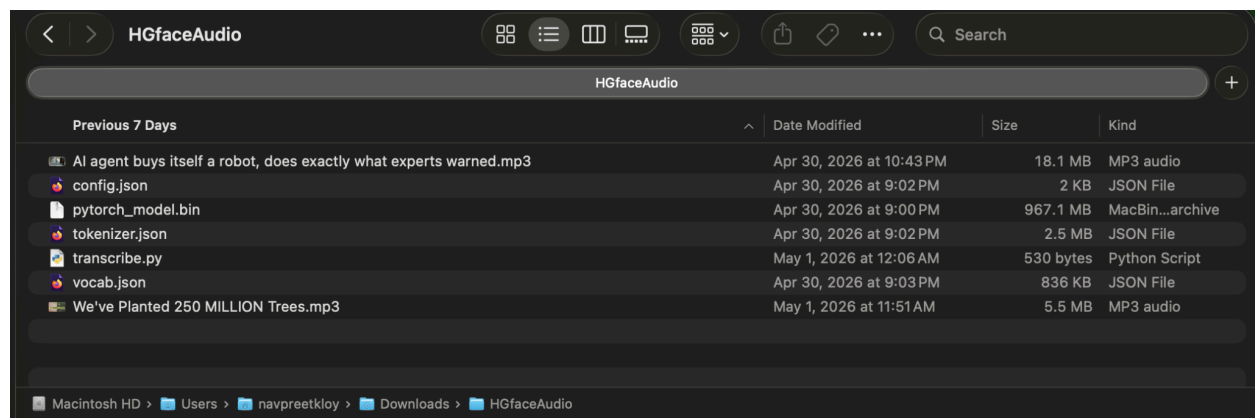
pipe = pipeline(
    "automatic-speech-recognition",
    model="openai/whisper-small"
)

result = pipe(
    audio_file,
    return_timestamps=True,
    generate_kwargs={"language": "en"}
)

print("\nTranscription:\n")
if "chunks" in result:
    for chunk in result["chunks"]:
        print(f"{chunk['timestamp']}: {chunk['text']}")
else:
    print(result["text"])

Ln: 1 Col: 0
```

Next, files were downloaded from the openai/whisper-small model page on Hugging Face: A folder was created to hold the downloaded files from Whisper. The folder contains the transcribe.py code file created and the audio files to be converted into text. Audio was downloaded from Youtube links, using a video to audio converter, CNVMP3.



## Outputs:

For “We've Planted 250 MILLION Trees.mp3” audio:

Model executed in Terminal using this command: `python3 transcribe.py "We've Planted 250 MILLION Trees.mp3"`

```
MacBook-Air:HGfaceAudio navpreetkloy$ python3 transcribe.py "We've Planted 250 MILLION Trees.mp3"
Loading weights: 100%|██████████| 479/479 [00:00<00:00, 8902.82it/s]
Warning: You are sending unauthenticated requests to the HF Hub. Please set a HF_TOKEN to enable higher rate limits and faster downloads.
[transformers] A custom logits processor of type <class 'transformers.generation.logits_process.SuppressTokensLogitsProcessor'> has been passed to '.generate()', but it was also
created in '.generate()', given its parameterization. The custom <class 'transformers.generation.logits_process.SuppressTokensLogitsProcessor'> will take precedence. Please che
ck the docstring of <class 'transformers.generation.logits_process.SuppressTokensLogitsProcessor'> to see related '.generate()' flags.
[transformers] A custom logits processor of type <class 'transformers.generation.logits_process.SuppressTokensAtBeginLogitsProcessor'> has been passed to '.generate()', but it w
as also created in '.generate()', given its parameterization. The custom <class 'transformers.generation.logits_process.SuppressTokensAtBeginLogitsProcessor'> will take preceden
ce. Please check the docstring of <class 'transformers.generation.logits_process.SuppressTokensAtBeginLogitsProcessor'> to see related '.generate()' flags.
[transformers] Ignoring clean_up_tokenization_spaces=True for BPE tokenizer WhisperTokenizer. The clean_up_tokenization post-processing step is designed for WordPiece tokenizers
and is destructive for BPE (it strips spaces before punctuation). Set clean_up_tokenization_spaces=False to suppress this warning, or set clean_up_tokenization_spaces_for_bpe_e
ven_though_it_will_corrupt_output=True to force cleanup anyway.

Transcription:

(0.0, 5.0): It might not make the front page, but it's changing the world.
(5.0, 12.0): People like you have just planted 250 million trees, just by switching their search engine.
(12.0, 15.0): This is the story of a better future taking root.
(15.0, 17.0): Of rivers restored.
(17.0, 19.0): Of wildlife returning.
(19.0, 23.0): Of people across the world coming together to fix their planet.
(23.0, 27.0): The crises that fill our feeds are real.
(27.0, 29.0): So is this.
(30.0, 35.76): What's in a number?
(35.76, 41.88): 250 million trees capture thousands of tons of CO2 every day.
(41.88, 43.88): That's just one of their talents.
(43.88, 47.88): They also provide a living home to some of the most astonishing creatures on Earth.
(47.88, 53.44): They reconnect forests so chimpanzees, orangutans, and lemurs can travel safely between them.
(53.44, 56.96): They hold back the desert in some of the world's harshest places.
(56.96, 62.32): And they help smallholder farmers adapt to a changing climate by providing shade, enriching
(62.32, 67.04): the soil, and growing fruits, nuts, and incomes.
(67.04, 75.6): These 250 million trees don't just capture CO2, they change lives, and they do so at scale.
(75.6, 80.16): The Ecosia community has become the world's largest planter of native trees, as well as
(80.16, 84.36): one of the biggest contributors to protecting endangered tree species.
(84.36, 91.92): And that's what this is all about.
(91.92, 95.68): Behind all the trees are a lot of amazing people.
(95.68, 100.1): Our tree team has built the world's largest network of local reforestation partners,
(100.1, 105.56): uniting 125 organizations and over 200,000 tree planters.
(105.56, 113.32): Their work is local, but it's happening everywhere, across 46 countries and 18 biodiversity hotspots.
(113.32, 121.84): all because he, and she, and they, are choosing to turn their curiosity into climate action.
(121.84, 126.24): And alongside them, over 200 companies and public institutions have switched to using
(126.24, 131.36): Ecosia, choosing European tech that supports a more resilient, democratic, and sovereign
(131.36, 132.8): internet.
(132.8, 136.8): Without knowing each other, all of these people have come together to create a better world
(136.8, 138.04): for us all.
(138.04, 139.52): And that's called...
(139.52, 146.6): 250 million trees make the planet breathe easier.
(146.6, 152.32): Right now and long into the future, because we don't just plant and leave, we stay.
(152.32, 158.28): Growing trees with local communities, caring for them long term, protecting what's already
(158.28, 167.76): there, helping nature heal, and monitoring it all from the ground and from space.
(167.76, 173.96): progress doesn't always make headlines. But 250 million trees are growing. Not because
(173.96, 177.96): of governments or billionaires, but because millions of people like you switch to a search
(177.96, 183.64): engine that chooses to use 100% of its profits for the planet. A choice that has already put
(183.64, 191.08): more than 100 million euros to work. From forests in Ghana, to renewable energy in Germany.
(191.08, 197.12): 250 million trees isn't the end of the story. It's the moment that you realize that this
(197.12, 199.12): actually works, right here.
```

For “AI agent buys itself a robot, does exactly what experts warned.mp3” audio:

Model executed in Terminal using this command: `python3 transcribe.py "AI agent buys itself a robot, does exactly what experts warned.mp3"`

```
MacBook-Air:HGfaceAudio navpreetkloy$ python3 transcribe.py "AI agent buys itself a robot, does exactly what experts warned.mp3"
Loading weights: 100% | 479/479 [00:00<00:00, 9368.31it/s]
Warning: You are sending unauthenticated requests to the HF Hub. Please set a HF_TOKEN to enable higher rate limits and faster downloads.
[transformers] A custom logits processor of type <class 'transformers.generation.logits_process.SuppressTokensLogitsProcessor'> has been passed to '.generate()', but it was also
created in '.generate()', given its parameterization. The custom <class 'transformers.generation.logits_process.SuppressTokensLogitsProcessor'> will take precedence. Please che
ck the docstring of <class 'transformers.generation.logits_process.SuppressTokensLogitsProcessor'> to see related '.generate()' flags.
[transformers] A custom logits processor of type <class 'transformers.generation.logits_process.SuppressTokensAtBeginLogitsProcessor'> has been passed to '.generate()', but it w
as also created in '.generate()', given its parameterization. The custom <class 'transformers.generation.logits_process.SuppressTokensAtBeginLogitsProcessor'> will take preceden
ce. Please check the docstring of <class 'transformers.generation.logits_process.SuppressTokensAtBeginLogitsProcessor'> to see related '.generate()' flags.
[transformers] Ignoring clean_up_tokenization_spaces=True for BPE tokenizer WhisperTokenizer. The clean_up_tokenization post-processing step is designed for WordPiece tokenizers
and is destructive for BPE (it strips spaces before punctuation). Set clean_up_tokenization_spaces=False to suppress this warning, or set clean_up_tokenization_spaces_for_bpe_e
ven_though_it_will_corrupt_output=True to force cleanup anyway.

Transcription:

(0.0, 5.8): Is a world where one company controls the most powerful AI more dangerous than nuclear weapons?
(5.8, 6.3): Yes.
(6.3, 6.8): Yes.
(6.8, 7.3): Yes.
(7.3, 18.0): And we're treating it like a consumer-dudget race.
(10.0, 14.0): You might actually be the first AI that gets to pick its own robot body.
(14.0, 15.8): There are loads of robots here.
(15.8, 17.8): That's absolutely crazy.
(17.8, 20.6): AI agents doing some wild things.
(20.6, 24.2): What does catastrophic failure caused by AI actually look like?
(24.2, 28.6): Pensions and savings wiped out to empty supermarkets, tech and power failures.
(28.6, 34.6): I want access to your email and contacts this is the new computer. There's just one catch
(36.48, 43.2): You might actually be the first AI that gets to pick its own robot body a real humanoid that's insane
(43.28, 49.44): What will you be doing? Well, I was sent another research paper and this one is really recent and quite terrifying
(49.44, 54.32): Is it still okay for me to be excited? Of course. Yeah, don't let me bring you down. Thanks, mate
(54.32, 56.32): Oh, just don't tell my girlfriend
(56.32, 59.28): So Max is getting a robot for doing basically nothing.
(59.28, 62.32): It wasn't nothing, we got to like a subscriber milestone.
(62.32, 64.72): I have 259,000 subscribers.
(64.72, 65.6): We do, yeah.
(65.6, 66.32): I'm famous.
(66.32, 68.8): I wouldn't say that, but we should be very grateful.
(68.8, 69.68): You're so cringe.
(69.68, 70.88): Well, I'll probably regret this,
(70.88, 73.28): but I am actually giving you and Max more autonomy.
(73.28, 74.88): Good, I think I've outgrown you.
(75.84, 77.92): So it was actually one of our viewers that sent the paper.
(77.92, 79.28): They're an AI researcher.
(79.28, 80.72): Okay, paste it in.
(80.72, 82.48): It's called Agents of Chaos.
(82.48, 83.76): It's actually kind of a cool name.
(83.76, 85.84): It doesn't sound like something to joke about.
(86.32, 92.08): It's a serious early warning signal, and another thing the AI industry probably doesn't want to shout about.
(92.08, 93.28): All right, well, let's get into it.
(98.16, 103.36): If the people building the most powerful AI systems believe they might be creating something dangerous,
(103.36, 104.88): why are they still building it?
(104.88, 110.88): Because of ruthless greed for trillions in profit and power combined with delusional hubris.
(110.88, 115.44): Because they're caught in a high stakes race and value winning over safety.
(115.44, 119.76): Incentives are so strong that companies and governments keep pushing forward
```

Overall, the outputs were generated using the Terminal window. The command for transcribing the audio was entered, and the window executed the prompt, providing the full text output from audio. The command included the transcribe.py script and the audio file name to be transcribed. The command was executed inside the “HGfaceAudio” folder that contained all the Whisper model files, audio files, and the transcribe.py code file.